



Language Modeling

Introduction to N-grams



Advanced smoothing algorithms

- Intuition used by many smoothing algorithms
 - Good--Turing
 - Kneser--Ney
 - Witten--Bell
- Use the count of things we've **seen once**
 - to help estimate the count of things we've **never seen**



Notation: N_c = Frequency of frequency c

N_c = the count of things we've seen c times

Sam I am I am Sam I do not eat

I	3
sam	2
am	2
do	1
not	1
eat	1

$$N_1 = 3$$

$$N_2 = 2$$

$$N_3 = 1$$

Good-Turing smoothing intuition



- You are fishing (a scenario from Josh Goodman), and caught:
- 10 carp, 3 perch, 2 whitefish, 1 trout, 1 salmon, 1 eel = 18 fish
 - How likely is it that next species is trout?
 - $1/18$
 - How likely is it that next species is new (i.e. catfish or bass)
 - Let's use our estimate of things-we-saw-once to estimate the new things.
 - $3/18$ (because $N_1=3$)
 - Assuming so, how likely is it that next species is trout?
 - Must be less than $1/18$
 - How to estimate?



Good Turing calculations

$$P_{GT}^*(\text{things with zero frequency}) = \frac{N_1}{N} \qquad c^* = \frac{(c+1)N_{c+1}}{N_c}$$

- Unseen (bass or catfish)
 - $c = 0$:
 - MLE $p = 0/18 = 0$
 - $P_{GT}^*(\text{unseen}) = N_1/N = 3/18$
- Seen once (trout)
 - $c = 1$
 - MLE $p = 1/18$
 - $C^*(\text{trout}) = 2 * N_2/N_1$
 $= 2 * 1/3$
 $= 2/3$
 - $P_{GT}^*(\text{trout}) = c^*/N = 2/3 / 18 = 1/27$



Resulting Good-Turing numbers

- Numbers from Church and Gale (1991)
- 22 million words of AP Newswire

$$c^* = \frac{(c+1)N_{c+1}}{N_c}$$

- It sure looks like $c^* = (c - .75)$

Count c	Good Turing c*
0	.0000270
1	0.446
2	1.26
3	2.24
4	3.24
5	4.22
6	5.19
7	6.21
8	7.24
9	8.25

Ney et al.'s Good Turing Intuition

H. Ney, U. Essen, and R. Kneser, 1995. On the estimation of 'small' probabilities by leaving-one-out.
IEEE Trans. PAMI. 17:12,1202--1212



Held-out words:

Ney et al. Good Turing Intuition (slide from Dan Klein)




Intuition from leave-one-out validation


- Take each of the c training words out in turn
- c training sets of size $c-1$, held-out of size 1
- What fraction of held-out words are unseen in training?
 - N_1/c (that occurred once)
- What fraction of held-out words are seen k times in training?
 - $(k+1)N_{k+1}/c$
- So in the future we expect $(k+1)N_{k+1}/c$ of the words to be those with training count k
- There are N_k words with training count k
- Each should occur with probability:
 - $(k+1)N_{k+1}/c/N_k$
- ...or expected count:

$$k^* = \frac{(k+1)N_{k+1}}{N_k}$$

Training



Held out





Good-Turing complications

(slide from Dan Klein)



- Problem: what about “the”? (say $c=4417$)
 - For small k , $N_k > N_{k+1}$
 - For large k , too jumpy, zeros wreck estimates
- Simple Good–Turing [Gale and Sampson]:
replace empirical N_k with a best-fit power law
once counts get unreliable
- The best-fit power law is a model that
predicts the probability of an event based on
its frequency, and it is chosen to fit the
observed counts as closely as possible.





Resulting Good-Turing numbers

- Numbers from Church and Gale (1991)
- 22 million words of AP Newswire

$$c^* = \frac{(c+1)N_{c+1}}{N_c}$$

Count c	Good Turing c*
0	.0000270
1	0.446
2	1.26
3	2.24
4	3.24
5	4.22
6	5.19
7	6.21
8	7.24
9	8.25



Absolute Discounting

- Save ourselves some time and just subtract 0.75 (or some d)!

discounted bigram

$$P_{\text{AbsoluteDiscounting}}(w_i | w_{i-1}) = \max\left(\frac{c(w_{i-1}, w_i) - d}{c(w_{i-1})}, 0\right)$$

- (Maybe keeping a couple extra values of d for counts 1 and 2)
- $d=0.75$ for $c>2$, smaller d for $c<2$
- $P(\text{Francisco} | \text{san}) = (c(\text{san francisco}) - 0.75) / c(\text{san})$



Absolute Discounting Interpolation

- Save ourselves some time and just subtract 0.75 (or some d)!

$$P_{\text{AbsoluteDiscounting}}(w_i | w_{i-1}) = \frac{\text{discounted bigram } c(w_{i-1}, w_i) - d}{c(w_{i-1})} + \lambda \overset{\text{Interpolation weight}}{(\overset{\swarrow}{w}_{i-1})} \underset{\nwarrow \text{unigram}}{P(w)}$$

- But should we really just use the regular unigram $P(w)$?
- Use continuation probability
 - Because of discounting you can estimate lambda directly don't need held out dataset for it
 - $\text{Lambda}(\text{san}) = d * \text{num of distinct words following san} / c(\text{san})$



Kneser-Ney Smoothing I

Better estimate for probabilities of lower-order unigrams!

- Shannon game: *I can't see without my reading—Francisco—?*
- “Francisco” is more common than “glasses”
- ... but “Francisco” always follows “San”
- The unigram is useful exactly when we haven't seen this bigram!
- Instead of $P(w)$: “How likely is w ”
- $P_{\text{continuation}}(w)$: “How likely is w to appear as a novel continuation?”
 - For each word, count the number of bigram types it completes
 - Every bigram type was a novel continuation the first time it was seen

$$P_{\text{CONTINUATION}}(w) \propto \left| \{w_{i-1} : c(w_{i-1}, w) > 0\} \right|$$



Kneser-Ney Smoothing II

How many times does w appear as a novel continuation:

$$P_{CONTINUATION}(w) \propto |\{w_{i-1} : c(w_{i-1}, w) > 0\}|$$

- Normalized by the total number of word bigram types

$$|\{(w_{j-1}, w_j) : c(w_{j-1}, w_j) > 0\}|$$

$$P_{CONTINUATION}(w) = \frac{|\{w_{i-1} : c(w_{i-1}, w) > 0\}|}{|\{(w_{j-1}, w_j) : c(w_{j-1}, w_j) > 0\}|}$$

Francisco is a continuation
 of how many unique
 words?

Count of all non-zero
 bigrams